



**ALLIANCE
UNIVERSITY**
Private University established in Karnataka State by Act No. 34 of year 2010
Recognized by the University Grants Commission (UGC), New Delhi.

**NAAC
GRADE A+**
ACCREDITED UNIVERSITY

Bachelor of Computer Applications

Semester – V

Deep Learning

Assignment No: 4

Title: Ethical Considerations in Deep Learning

Submitted By

Name: Dipin Roka

Reg. No.:2411021240007

Date:06/09/2026

[Githublink](#)

Submitted To:

Faculty Name: Nagasri p

Department of Computer Application
Alliance University
Chandapura - Anekal Main Road, Anekal
Bengaluru – 562106

Project 4: Ethical Considerations in Deep Learning

Description

Students will explore ethical challenges associated with Deep Learning systems, including bias, privacy, fairness, transparency, and responsible use of AI. This report addresses three major ethical issues, examines a real-world incident, and recommends practical measures to promote responsible development and deployment of deep learning models.

Introduction

Deep learning models now support decisions in recruitment, banking, healthcare, policing, and education. Unlike traditional software with fixed rules, these models learn behaviour from large datasets and adapt their internal logic during training. This learning ability gives them great predictive power, but it also creates new ethical risks. When data is unbalanced, when personal information is collected without clear permission, or when a model's reasoning cannot be explained, the result can be unfair or unsafe outcomes. This report discusses three such concerns, illustrates them through a documented industry case, and outlines steps that developers and organisations can take to reduce harm.

Fairness Privacy Transparency Human Oversight

Task 1: Three Ethical Issues in Deep Learning

1. Bias and Unfair Treatment

A deep learning model reflects the composition of the data it is trained on. If certain communities, genders, or regions are under-represented, the model tends to produce lower accuracy or harsher judgments for those groups. Because the model only optimises for overall error reduction, it does not automatically correct for social inequality. In practice this has meant resume-screening systems favouring male candidates, skin-tone dependent errors in face analysis, and loan models assigning higher risk to already disadvantaged applicants. The harm is cumulative: an unfair model deployed at scale can deny opportunities to thousands of people while appearing objective.




Figure 1: Illustrative gap in shortlist rates before vs. after bias mitigation.

2. Privacy and Consent in Large-Scale Data Collection

Modern neural networks improve with more data, which creates pressure to gather text, images, voice, and location records in bulk. Much of this material is taken from public websites, applications, and surveillance sources where individuals never gave informed consent for AI training. Once personal details are embedded in model weights, they are difficult to remove. Researchers have also shown that attackers can sometimes reconstruct faces, text fragments, or medical details from a released model. Ethical use therefore requires asking not only whether data is available, but whether its collection respected user autonomy and legal limits.

3. Opacity and Difficulty in Assigning Responsibility

Deep networks contain millions of parameters distributed across many layers, so even their creators cannot always trace why a single prediction was made. This lack of interpretability becomes critical when a model recommends rejecting a job applicant, denying bail, or altering a medical diagnosis. Affected persons cannot appeal effectively if no reason is given, and managers may place blind trust in a score they do not understand. Without clear logs, explanations, and defined ownership, it becomes unclear whether the data team, the deploying company, or the end user should answer for a mistake.

Task 2: Real-World Example — Amazon's Automated Hiring System

Between 2014 and 2018, Amazon developed a deep learning tool to review resumes and shortlist candidates for technical roles. The system was trained on ten years of past applications received by the company, most of which came from men.

During testing, the tool learned to downgrade resumes that contained the word “women’s”, such as “women’s chess club captain”, and to prefer language more common in male applicants’ resumes. Even after attempts to remove explicit gender indicators, the model found indirect proxies and continued to produce gender-skewed rankings. Amazon ultimately discontinued the project and never deployed it for final hiring decisions.

Table 1: Timeline of the Amazon hiring-tool case

Period	What happened	Ethical signal
2014–2017	Model trained on 10 years of male-dominated resumes	Historical bias enters data
2018 (testing)	System penalises resumes with women’s-related terms	Bias detected in audit
2018 (later)	Project scrapped after fixes fail	Human oversight stops deployment

Why this matters ethically:

- **Fairness:** The model converted a historical gender imbalance into an automated selection rule, penalising qualified women.
- **Transparency:** Candidates had no knowledge that an algorithm was shaping their chances, and no way to contest its logic.
- **Responsibility:** The incident showed that high overall accuracy on past data does not guarantee equitable behaviour in the future, and that human review must remain mandatory for employment decisions.

Task 3: Measures to Reduce the Identified Issues

Responsible Deep Learning Lifecycle



Figure 2: Responsible deep-learning lifecycle — from consented data collection to continuous monitoring.

- **Balance and audit training data:** Check representation across gender, age, region, and language before training, and supplement or re-weight samples where gaps are found.

- Apply fairness testing, not only accuracy testing: Measure false-positive and false-negative rates separately for each group and set minimum fairness thresholds for release.
- Strengthen data governance: Collect only necessary data, obtain clear consent, anonymise sensitive fields, and document data sources so removal requests can be honoured.
- Add explainability: Provide reason codes or feature-importance summaries for high-impact decisions, using tools and design practices suited to non-technical users.
- Keep humans in control: Use model outputs as suggestions. A trained reviewer should make the final call in hiring, lending, medical, and legal contexts, with authority to override the system.
- Require external checks: Invite independent audits for high-risk systems and maintain versioned records of data, code, and evaluation results for inspection.
- Monitor after deployment: Track live performance by group, collect appeals and complaints, and retrain or suspend the model if disparity increases over time.

Conclusion

Deep learning can improve speed and consistency, but without ethical safeguards it can also reproduce past discrimination, erode privacy, and hide accountability behind technical complexity. The Amazon hiring case demonstrates that problems often originate in training data and remain invisible until specifically tested. By combining representative data, fairness evaluation, privacy-respecting collection, interpretable outputs, human oversight, and continuous monitoring, organisations can use deep learning in a way that is both effective and responsible.